



# Training DNNs in Tiny Subspaces

Introductory Talk | Janik Philipps

# **Low Dimensional Landscape Hypothesis is True: DNNs can be Trained in Tiny Subspaces**

Tao Li<sup>1</sup>      Lei Tan<sup>1</sup>      Qinghua Tao<sup>2</sup>      Yipeng Liu<sup>3</sup>      Xiaolin Huang<sup>\*1</sup>  
<sup>1</sup>Shanghai Jiao Tong University      <sup>2</sup>Tsinghua University  
<sup>3</sup>University of Electronic Science and Technology

# Motivation: Deep neural networks depend on million of parameters causing severe problems

---

Models	# Parameter
VGG11_bn [41]	28.5M
EfficientNet-B0 [47]	4.14M
MobileNet [20]	3.3M
DenseNet121 [22]	7.0M
Inceptionv3 [46]	22.3M
Xception [7]	21.0M
GoogLeNet [45]	6.2M
ShuffleNetv2 [35]	1.3M
SqueezeNet [23]	0.78M
SEResNet18 [21]	11.4M
NasNet [53]	5.2M

- ▶ Potential overfitting of the data
- ▶ Only first order methods applicable
- ▶ Computationally and time intensive training
- ▶ Many training data needed
- ▶ Questionable solution quality

## Motivation: The number of independent optimization variables may be smaller than the number of parameters

---

- ▶ Due to strong mutual relationships, **regarding each parameter** of deep neural networks **as an independent variable is too rough**
- ▶ The **gradients of parameters are strongly related** due to the training via backpropagation
- ▶ The parameters in the same layers also have **synergy correlations**



"DNNs can be trained in low-dimensional subspaces"

## Motivation: Pioneering work achieved 90% accuracy in smaller spaces

---


- ▶ There exists pioneering work on **training via random projections**
- ▶ For example, on CIFAR-10, LeNet with 62006 parameters could be optimized in 2900-dimensional subspaces with **90% accuracy of regular SGD training**.



Very promising but we can do even better!

Many standard neural network architectures could be **well trained by only 40 independent variables** with almost the same performance.

## Approach in the paper




- ▶ There are **three parameters**  $w(1)$ ,  $w(2)$ ,  $w(3)$  to **optimize**.
- ▶ The training trajectory  $\{w_i\}_{i=0,\dots,t}$  could be in a **two-dimensional subspace** spanned by  $e_1$  and  $e_2$
- ▶ Training in the low-dimensional space can have **comparable performance** as training in the regular space



# Agenda

---



- ▶ Dynamic Linear Dimensionality Reduction (DLDR)
- ▶ DLDR-based Quasi-Newton Algorithm
- ▶ Numerical Experiments

# Agenda

---



- ▶ **Dynamic Linear Dimensionality Reduction (DLDR)**
- ▶ DLDR-based Quasi-Newton Algorithm
- ▶ Numerical Experiments



# Dynamic Linear Dimensionality Reduction

---

Assumption: The layer width is unlimited.

- ▶ Under infinite-width limit, a wide neural network estimator can be approximated by a linear model under gradient descent, such that

$$f^{lin}(x, w_t) \approx f(x, w_0) + \nabla_w f(\mathcal{X}, w_0)(w_t - w_0)$$

- ▶ We formulate the gradient flow of a single-output neural network:

$$\dot{w}_t = -\nabla_w f(\mathcal{X}, w_t)^T \nabla_{f_t(\mathcal{X}, w_t)} \mathcal{L}$$

- ▶ Using the linearized model, the dynamics of gradient flow are governed by

$$\dot{w}_t = -\nabla_w f(\mathcal{X}, w_0)^T \nabla_{f^{lin}(\mathcal{X}, w_t)} \mathcal{L}$$

# Dynamic Linear Dimensionality Reduction

---

- ▶ Applying Singular Value Decomposition on  $\nabla_w f(\mathcal{X}, w_0)$  yields

$$\nabla_w f(\mathcal{X}, w_0) = U_0 \Sigma_0 V_0^T \in \mathbb{R}^{m \times n}$$

- ▶ By definition of the Neural Tangent Kernel, we have

$$\Theta_0 = \nabla_w f(\mathcal{X}, w_0) \nabla_w f(\mathcal{X}, w_0)^T = U_0 \Sigma_0 \Sigma_0^T U_0^T$$

- ▶ Based on the properties of NTK in infite-width limit, we can approximate  $\Sigma_0$  by a low-rank matrix  $\tilde{\Sigma}_0 \in \mathbb{R}^{d \times d}$  containing the first  $d$  largest singular values, such that

$$\Sigma_0 \approx \tilde{U}_0 \tilde{\Sigma}_0 \tilde{V}_0^T$$

- ▶ Thus we derive the approximations

$$\begin{aligned} \nabla_w f(\mathcal{X}, w_0) &\approx U_0 \tilde{U}_0 \tilde{\Sigma}_0 \tilde{V}_0^T V_0^T \\ \dot{w}_t &\approx -V_0 \tilde{V}_0 (\tilde{\Sigma}_0 \tilde{U}_0^T U_0^T \nabla_{f^{lin}(\mathcal{X}, w_t)} \mathcal{L}) \end{aligned}$$

How to find the low-dimensional subspace?

- ▶ 1) Sample  $t$  steps of parameters during the training, namely,  $\{w_1, w_2, \dots, w_t\}$ .
- ▶ 2) Centralize these as  $\bar{w} = \frac{1}{t} \sum_{i=1}^t w_i$  and let  $W = [w_1 - \bar{w}, \dots, w_t - \bar{w}]$ .
- ▶ 3) Find a  $d$ -dimensional subspace spanned by  $P = [e_1, e_2, \dots, e_d]$  to cover  $W$ .

The third step is to find a subspace that the distance of  $W$  and  $P^T W$  is minimized.

- ▶ We consider the SVD  $W = U \Sigma V^T$  with  $U = [u_1, \dots, u_n]$ ,  $\Sigma = \text{diag}(\sigma_1, \dots, \sigma_t)$  and  $V = [v_1, \dots, v_t]$ . The first  $d$  columns of  $U$  are the independent variables.
- ▶ Compute  $v_i$  for  $i = 1, \dots, d$  by the spectral decomposition of  $W^T W$  and derive

$$W v_i = \sigma_i u_i, \quad i = 1, \dots, d.$$

---

## Algorithm 1 Dynamic Linear Dimensionality Reduction (DLDR)

---

- 1: Sample parameter trajectory  $\{\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_t\}$  along the training;
  - 2:  $\bar{\mathbf{w}} = \frac{1}{t} \sum_{i=1}^t \mathbf{w}_i$ ;
  - 3:  $W = [\mathbf{w}_1 - \bar{\mathbf{w}}, \mathbf{w}_2 - \bar{\mathbf{w}}, \dots, \mathbf{w}_t - \bar{\mathbf{w}}]$ ;
  - 4: Perform spectral decomposition on  $W^\top W$  and obtain the largest  $d$  eigenvalues  $[\sigma_1^2, \sigma_2^2, \dots, \sigma_d^2]$  with the corresponding eigenvectors  $[\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_d]$ ;
  - 5:  $\mathbf{u}_i = \frac{1}{\sigma_i} W \mathbf{v}_i$ ;
  - 6: Return  $[\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_d]$  as the orthonormal bases.
-

# Dynamic Linear Dimensionality Reduction

Setting: learning rate 0.1, batch size 128, SGD with momentum of 0.9, 3 trials total


dimension reduction method	dimension	test accuracy
—	78330	83.84 $\pm$ 0.34
Li <i>et al.</i> [31]	7982	58.35 $\pm$ 0.04
Gressmann <i>et al.</i> [15]	7982	70.26 $\pm$ 0.02
DLDR (ours)	<b>15</b>	<b>85.18<math>\pm</math>0.20</b>

Test accuracy of ResNet8 on CIFAR-10



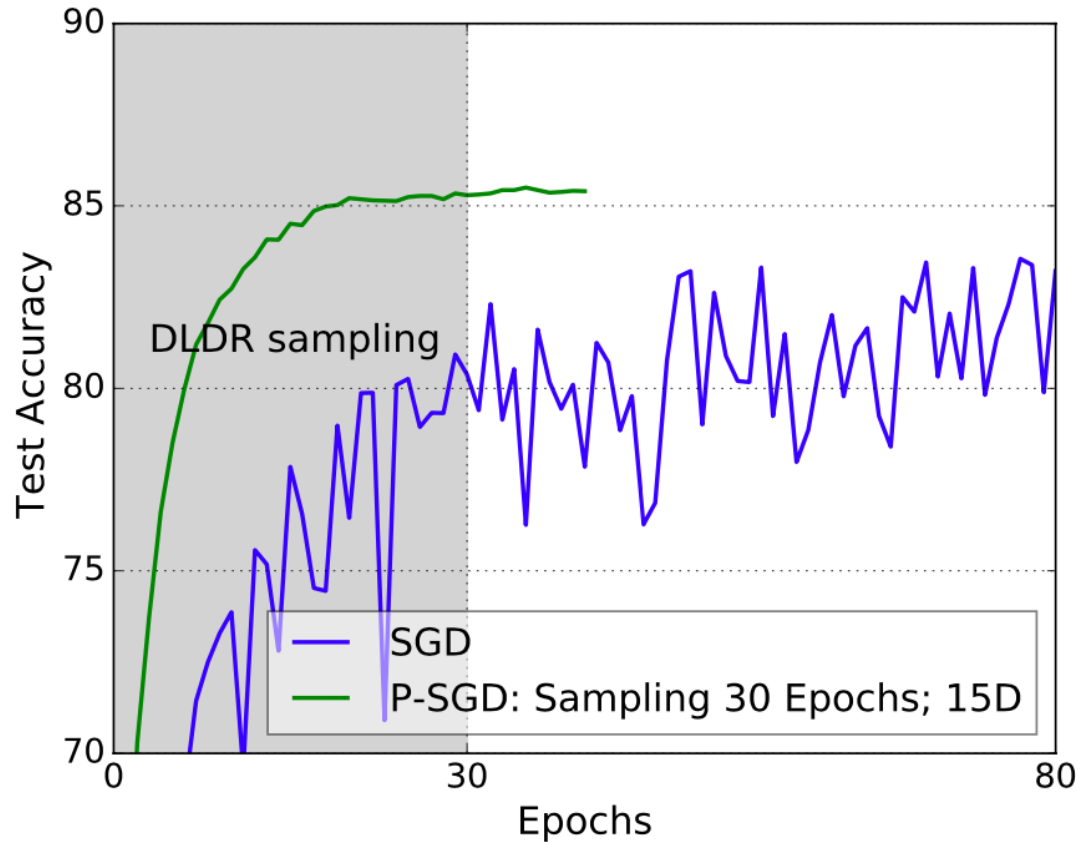
# Dynamic Linear Dimensionality Reduction

---



PCA ratio of the first 30 epochs training trajectory in the principal directions

# Dynamic Linear Dimensionality Reduction



Training performance of SGD and P-SGD (in 15D subspace)

# Agenda

---



- ▶ Dynamic Linear Dimensionality Reduction (DLDR)
- ▶ **DLDR-based Quasi-Newton Algorithm**
- ▶ Numerical Experiments

# DLDR-based Quasi-Newton Algorithm

---

- ▶ Working in the low-dimensional space has the advantage that we can approximate the Hessian matrix as

$$H \approx PH_0P^T$$

- ▶ We use the BFGS algorithm with rank two correction update to approximate the inverse Hessian matrix as

$$B_{k+1} = V_k^T B_k V_k + \rho_k \tilde{s}_k \tilde{s}_k^T$$

$$y_k = \tilde{g}_{k+1} - \tilde{g}_k$$

$$\rho_k = (y_k^T \tilde{s}_k)^{-1}$$

$$V_k = I - \rho_k y_k \tilde{s}_k^T$$

where  $\tilde{g}_k = P^T g_k$  is the projected gradient and  $\tilde{s}_k = P^T(w_{k+1} - w_k)$  is the projected difference between the parameters

---

## Algorithm 2 P-BFGS

---

- 1: Obtain  $P = [\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_d]$  by DLDR in Algorithm 1;
  - 2: Initialize  $k \leftarrow 0$  and  $B_k \leftarrow I$ ;
  - 3: **while** not converging **do**
  - 4:     Sample the mini-batch data  $\mathcal{B}_k$ ;
  - 5:     Compute the gradient  $\mathbf{g}_k$  on  $\mathcal{B}_k$ ;
  - 6:     Perform the projection  $\tilde{\mathbf{g}}_k \leftarrow P^\top \mathbf{g}_k$ ;
  - 7:     **if**  $k > 0$  **then**
  - 8:         Do Quasi-Newton update  $B_k$  with (13);
  - 9:     Compute  $\alpha_k$  using backtracking line search;
  - 10:      $\tilde{\mathbf{s}}_k \leftarrow -\alpha_k B_k \tilde{\mathbf{g}}_k$ ;
  - 11:      $\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k + P \tilde{\mathbf{s}}_k$ ;      $\triangleright$  Update the parameters
  - 12:      $k \leftarrow k + 1$ ;
-



# Agenda

---



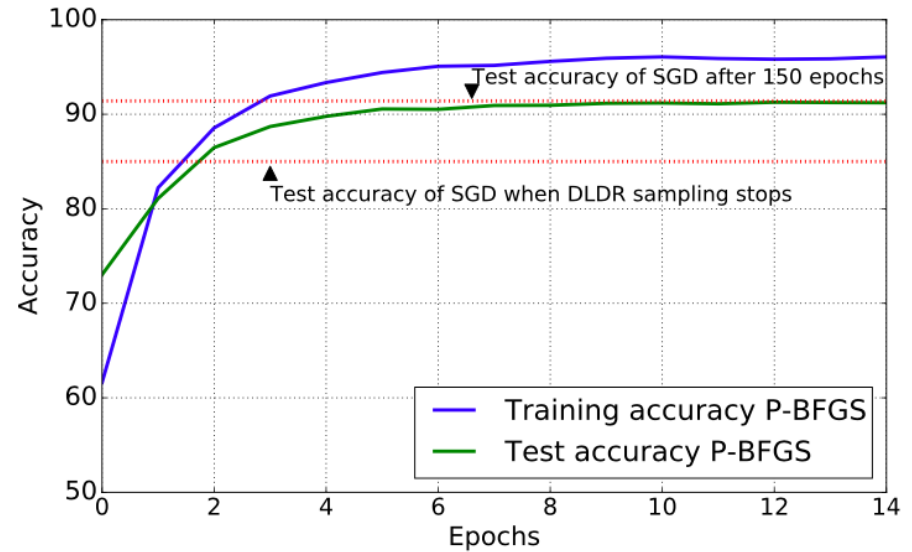
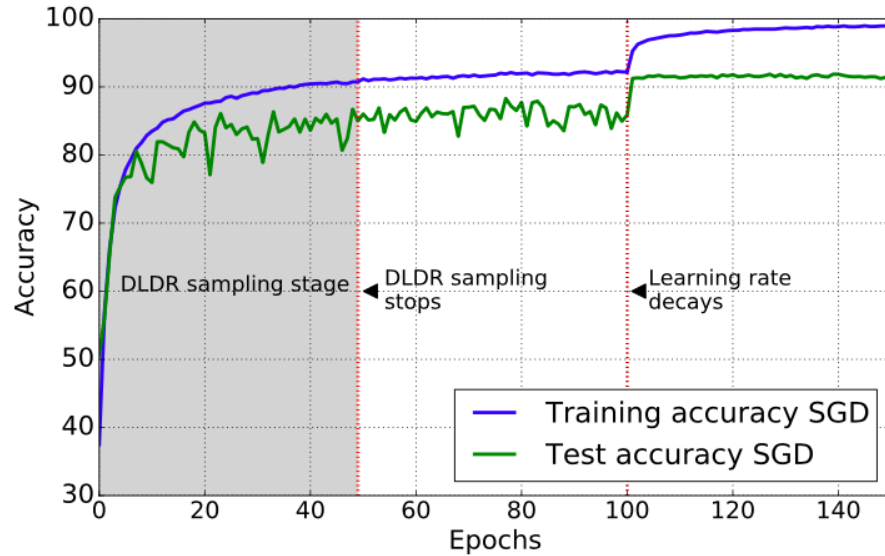
- ▶ Dynamic Linear Dimensionality Reduction (DLDR)
- ▶ DLDR-based Quasi-Newton Algorithm
- ▶ **Numerical Experiments**

## Numerical Experiments

Models	# Parameter	SGD (#training epochs)			P-SGD (#sample epochs)	
		50	100	200	50	100
VGG11_bn [41]	28.5M	58.38	59.90	68.87	68.72	70.18
EfficientNet-B0 [47]	4.14M	62.53	63.68	72.94	71.68	72.64
MobileNet [20]	3.3M	57.15	58.67	67.94	66.86	68.00
DenseNet121 [22]	7.0M	65.39	64.57	76.76	74.25	76.34
Inceptionv3 [46]	22.3M	61.68	64.00	76.25	75.15	76.83
Xception [7]	21.0M	64.57	65.81	75.47	75.68	75.56
GoogLeNet [45]	6.2M	62.32	66.32	76.88	75.66	77.27
ShuffleNetv2 [35]	1.3M	62.90	63.15	72.06	71.34	72.29
SqueezeNet [23]	0.78M	59.52	58.56	70.29	69.89	70.60
SEResNet18 [21]	11.4M	64.74	64.68	74.95	74.33	75.09
NasNet [53]	5.2M	63.73	66.80	77.34	77.19	77.03

Test accuracy on CIFAR-100 for regular training and optimization in 40D subspaces

# Numerical Experiments




The accuracy of ResNet20 on CIFAR-10 in different epochs

## Numerical Experiments

Dataset		CIFAR-10	CIFAR-100	ImageNet
Model		ResNet20	ResNet32	ResNet18
SGD	epochs	150	200	90
	acc	$91.55 \pm 0.23$	$68.40 \pm 0.45$	69.794
P-BFGS	sampling	80	100	60
	epochs	20	20	4
	acc	$91.72 \pm 0.10$	$69.90 \pm 0.48$	69.720

Test accuracy obtained by SGD and P-BFGS

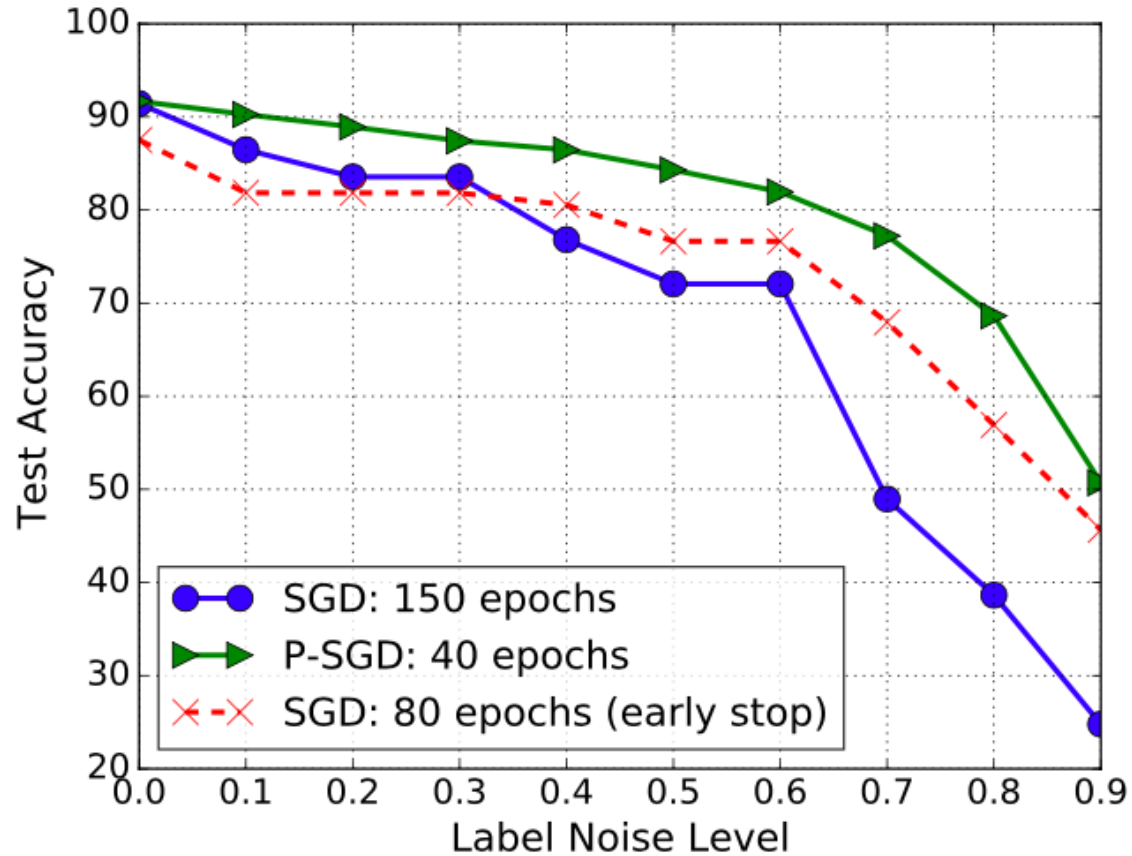
# Numerical Experiments



Wall-clock time consumption comparisons



# Numerical Experiments



The performance under different level of label noise

---

# Are there any questions?



## Ideas for further exploration

---

- ▶ Test different second order methods
- ▶ Improve the subspace extraction (different sampling strategies, different dimensions, etc.)

**Further ideas?**